

# Constraints for Regularization

---

Where there's room for one elephant, there's room for two

# Something Fishy is Going On

Notice how we are consistently getting 0 RSSE?

```
In [2]: rflows, rpaths = util.solve_path_selection_full(tug, node_counts, arc_counts, verbose=0, so
print('FLOW: PATH')
util.print_solution(tug, rflows, rpaths, sort='descending')
sse = util.get_reconstruction_error(tug, rflows, rpaths, node_counts, arc_counts)
print(f'\nRSSE: {np.sqrt(sse):.2f}')
```

```
FLOW: PATH
8.17: 2,3 > 3,3
5.47: 0,2 > 1,2 > 2,2 > 3,2
3.74: 3,3
3.10: 0,1 > 1,1 > 2,0 > 3,0
1.79: 1,0 > 2,0 > 3,0
1.79: 0,1 > 1,1 > 2,0 > 3,2
1.53: 1,0 > 2,0 > 3,2
```

```
RSSE: 0.00
```

# Something Fishy is Going On

Notice how we are consistently getting 0 RSSE?

```
In [2]: rflows, rpaths = util.solve_path_selection_full(tug, node_counts, arc_counts, verbose=0, so
print('FLOW: PATH')
util.print_solution(tug, rflows, rpaths, sort='descending')
sse = util.get_reconstruction_error(tug, rflows, rpaths, node_counts, arc_counts)
print(f'\nRSSE: {np.sqrt(sse):.2f}')
```

```
FLOW: PATH
8.17: 2,3 > 3,3
5.47: 0,2 > 1,2 > 2,2 > 3,2
3.74: 3,3
3.10: 0,1 > 1,1 > 2,0 > 3,0
1.79: 1,0 > 2,0 > 3,0
1.79: 0,1 > 1,1 > 2,0 > 3,2
1.53: 1,0 > 2,0 > 3,2
```

```
RSSE: 0.00
```

How can that be the case?

# The Must be Some Noise in Your Dataset

So far, we have implicitly assumed **noiseless data**

We will fix that by adding some **proportional noise**

- Which we picked since it is reasonably realistic
- ...Even if it causes issues for our MSE loss

**This is done in the `add_proportional_noise` function:**

```
# Add noise to the node counts
for k, v in node_counts.items():
    node_counts[k] = max(0, v * (1 + np.random.normal(0, sigma)))
# Add noise to the arc counts
for k, v in arc_counts.items():
    arc_counts[k] = max(0, v * (1 + np.random.normal(0, sigma)))
```

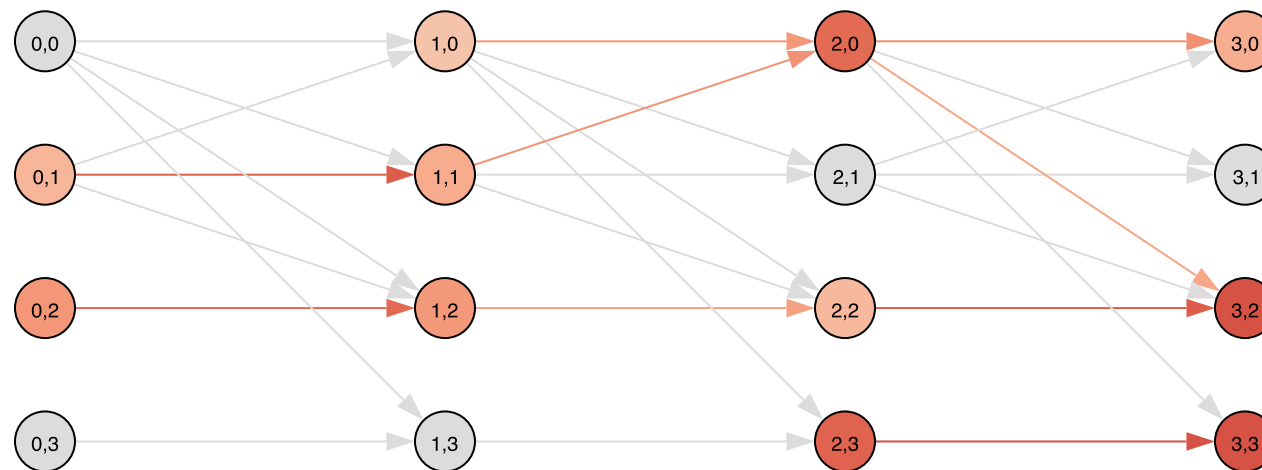
- The `sigma` parameter controls the noise level

# There Must be Some Noise in Your Dataset

Let's inject **a lot** of noise and inspect the results

```
In [3]: node_counts_n, arc_counts_n = util.add_proportional_noise(node_counts, arc_counts, sigma=0.2)
visual_style = util.get_visual_style(tug, vertex_weights=node_counts_n, edge_weights=arc_counts_n)
fig.plot(tug, **visual_style, bbox=(700, 300), margin=50)
```

Out[3]:



# Solving the Noisy Path Formulation

Let's try to solve the path formulation with noisy data

```
In [4]: rflows_n, rpaths_n = util.solve_path_selection_full(tug, node_counts_n, arc_counts_n, verbose=True)
print('FLOW: PATH')
util.print_solution(tug, rflows_n, rpaths_n, sort='descending', max_paths=15)
sse = util.get_reconstruction_error(tug, rflows_n, rpaths_n, node_counts_n, arc_counts_n)
print(f'RSSE: {np.sqrt(sse):.2f}')
```

```
FLOW: PATH
7.28: 2,3 > 3,3
4.49: 3,3
3.02: 0,2 > 1,2 > 2,2 > 3,2
2.22: 0,1 > 1,1 > 2,0 > 3,0
2.19: 0,2 > 1,2
2.08: 3,2
1.48: 1,0 > 2,0 > 3,2
1.47: 1,0 > 2,0 > 3,0
1.43: 2,3
1.28: 2,2 > 3,2
1.27: 0,1 > 1,1 > 2,0 > 3,2
0.90: 0,1 > 1,1
0.84: 0,2
0.60: 3,0
0.45: 1,0 > 2,0
...
RSSE: 1.64
```

# Solving the Noisy Path Formulation

There are some very noticeable differences w.r.t. the baseline

- The RSSE is a bit higher, which could be expected
- But there are also **many more** paths, and they tend to be **shorter**

**What is going on?**

# Solving the Noisy Path Formulation

**There some very noticeable differences w.r.t. the baseline**

- The RSSE is a bit higher, which could be expected
- But there are also **many more** paths, and they tend to be **shorter**

**What is going on?**

**We have overfitting issues**

- Our data-mining model is **almost free of bias** (we can use any possible path)
- Hence, the model tries to cover all nodes with many, short, paths

**Can we do something about it?**



# L1 Regularization, Put to Its Purpose

We know that we can use an L1 regularizer to encourage longer paths

...After all, L1 and L2 regularization were born to counter overfitting

```
In [5]: rflows_n2, rpaths_n2 = util.solve_path_selection_full(tug, node_counts_n, arc_counts_n, alpl
print('FLOW: PATH')
util.print_solution(tug, rflows_n2, rpaths_n2, sort='descending')
sse = util.get_reconstruction_error(tug, rflows_n2, rpaths_n2, node_counts_n, arc_counts_n)
print(f'RSSE: {np.sqrt(sse):.2f}')
```

```
FLOW: PATH
7.92: 2,3 > 3,3
4.80: 0,2 > 1,2 > 2,2 > 3,2
2.04: 0,1 > 1,1 > 2,0 > 3,0
1.87: 0,1 > 1,1 > 2,0 > 3,2
1.35: 1,0 > 2,0 > 3,2
1.35: 1,0 > 2,0 > 3,0
0.55: 0,1 > 1,1 > 2,0 > 3,3
0.46: 1,0 > 2,0 > 3,3
RSSE: 4.96
```

We get fewer, longer paths, at the expense of a higher RSSE

# L1 Regularization, Put to Its Purpose

As usual, we can try to improve our results via consolidation

```
In [6]: node_counts_r, arc_counts_r = util.get_counts(tug, rflows_n2, rpaths_n2)
        cflows, cpaths, cflag = util consolidate_paths(tug, rpaths_n2, node_counts_r, arc_counts_r)
        print('FLOW: PATH')
        util.print_solution(tug, cflows, cpaths, sort='descending', max_paths=5)
```

```
FLOW: PATH
7.92: 2,3 > 3,3
4.80: 0,2 > 1,2 > 2,2 > 3,2
3.39: 0,1 > 1,1 > 2,0 > 3,0
2.16: 1,0 > 2,0 > 3,2
1.07: 0,1 > 1,1 > 2,0 > 3,2
...
```

```
In [7]: util.print_ground_truth(flows, paths, sort='descending')
```

```
8.17: 2,3 > 3,3
5.47: 0,2 > 1,2 > 2,2 > 3,2
4.89: 0,1 > 1,1 > 2,0 > 3,0
3.74: 3,3
3.32: 1,0 > 2,0 > 3,2
```

- We got many of the ground truth paths correctly, but we are still using many spurious ones

# Minimum Cover Constraints

**Perhaps we could try to counter the adverse effects of the L1 term**

...Without losing all of its benefits

- If the L1 weight is too low, the regularizer has little effect
- ...But if it is too high, the solver stops focusing on the count reconstruction

**What can we do?**

# Minimum Cover Constraints

Perhaps we could try to counter the adverse effects of the L1 term

...Without losing all of its benefits

- If the L1 weight is too low, the regularizer has little effect
- ...But if it is too high, the solver stops focusing on the count reconstruction

What can we do?

One way to achieve this consists in introducing **new constraints**

- For example, we could require for each vertex
- ...To recover a **minimum fraction**  $\gamma$  of the total count, i.e.

$$\sum_v x_v \geq \gamma v$$

# Minimum Cover Constraints

The path formulation then becomes

$$\arg \min_x \left\{ \frac{1}{2} x^T P x + q^T x \mid V x \geq \gamma v, x \geq 0 \right\}$$

With  $P = V^T V + E^T E$  and  $q = -V^T v - E^T \hat{e} + \alpha$

- We have incorporated both the L1 term (them  $\alpha$  term)
- ...And the minimum cover constraints

## When calling the OSQP solver

- We need to include  $\alpha$  in the definition of  $q$
- Then we need to extend the constraint matrix/vectors  $A, l, u$
- ...So as to account for  $Vx \geq \gamma v$

# Solving the Modified Path Formulation

Let's try to solve the problem for  $\alpha = 3$  and  $\gamma = 0.8$

```
In [8]: rflows_n3, rpaths_n3 = util.solve_path_selection_full(tug, node_counts_n, arc_counts_n, alpl
                                                min_vertex_cover=0.95, solver='piqp')
print('FLOW: PATH')
util.print_solution(tug, rflows_n3, rpaths_n3, sort='descending', max_paths=10)
sse = util.get_reconstruction_error(tug, rflows_n3, rpaths_n3, node_counts_n, arc_counts_n)
print(f'RSSE: {np.sqrt(sse):.2f}')
```

```
FLOW: PATH
8.28: 2,3 > 3,3
5.23: 0,2 > 1,2 > 2,2 > 3,2
2.62: 3,3
2.42: 0,1 > 1,1 > 2,0 > 3,0
1.87: 0,1 > 1,1 > 2,0 > 3,2
1.78: 1,0 > 2,0 > 3,0
1.52: 1,0 > 2,0 > 3,2
0.63: 3,2
0.41: 0,2 > 1,2
0.14: 0,1 > 1,1 > 2,0 > 3,3
...
RSSE: 3.13
```

```
/Users/michelelombardi/Library/Caches/pypoetry/virtualenvs/mining-with-co-eC0Z9-8k-py3.11/
lib/python3.11/site-packages/scipy/sparse/_index.py:108: SparseEfficiencyWarning: Changing
the sparsity structure of a csc_matrix is expensive. lil_matrix is more efficient.
self._set_intXint(row, col, x.flat[0])
```

# Solving the Modified Path Formulation

Let's see what happens with consolidation

```
In [9]: node_counts_r, arc_counts_r = util.get_counts(tug, rflows_n3, rpaths_n3)
        cflows, cpaths, cflag = util.consolidate_paths(tug, rpaths_n3, node_counts_r, arc_counts_r)
        print('FLOW: PATH')
        util.print_solution(tug, cflows, cpaths, sort='descending', max_paths=5)
```

```
FLOW: PATH
8.28: 2,3 > 3,3
5.23: 0,2 > 1,2 > 2,2 > 3,2
4.28: 0,1 > 1,1 > 2,0 > 3,0
3.40: 1,0 > 2,0 > 3,2
2.62: 3,3
...
```

```
In [10]: util.print_ground_truth(flows, paths, sort='descending')
```

```
8.17: 2,3 > 3,3
5.47: 0,2 > 1,2 > 2,2 > 3,2
4.89: 0,1 > 1,1 > 2,0 > 3,0
3.74: 3,3
3.32: 1,0 > 2,0 > 3,2
```

We got all paths right, and the flows are **closer** to their real value

# Column Generation with Constraints in the Master

---

Where we make our first acquaintance with the KKT conditions



## CG and Modified Path Formulation

The new formulation requires to add new constraints **in the master**

This is a problem for Column Generation. Due to the constraints:

- Just looking at the gradient may now be misleading
- ...Since changing a variable may **force to change others**

$$\arg \min_x \left\{ \frac{1}{2} x^T P x + q^T x \mid V x \geq \gamma v, x \geq 0 \right\}$$

**We need a "constraint-aware gradient"**

One way to achieve that is to rely on a **Lagrangian approach**

- The idea is to turn the constraints in to **cost terms**
- ...And control their satisfaction by **adjusting weights (multipliers)**

We will discuss this approach in a general setting

# Lagrangian Approach

Let's consider an optimization in the form

$$\operatorname{argmin}_x \{ f(x) \mid g(x) \leq 0 \} \quad (\mathbf{P1})$$

where  $\mathbf{x}$  belongs to  $\mathbb{R}^n$  (i.e. this is numeric optimization)

**From this, we can obtain a related, unconstrained optimization problem**

...By moving the constraints in to the cost function, with **weights/multipliers  $\lambda$** :

$$\operatorname{argmin}_x \mathcal{L}(x, \lambda) = f(x) + \lambda^T g(x) \quad (\mathbf{P2})$$

The term  $\mathcal{L}(x, \lambda)$  is called a **Lagrangian**

- If a constraint  $g_i(\mathbf{x})$  is violated,  $\mathcal{L}$  gets a penalty w.r.t.  $f(\mathbf{x})$
- If a constraint  $g_i(\mathbf{x})$  is satisfied,  $\mathcal{L}$  gets a reward w.r.t.  $f(\mathbf{x})$

**We want to solve (P1) by controlling the multipliers in (P2)**

## ...And KKT Conditions

Let's assume that  $x$  is an **local** optimum for the original problem

...If we want to reach it via (P2), the multipliers should be **just right**:

- They should make the Lagrangian gradient null, i.e.

$$\nabla_x \mathcal{L}(x, \lambda) = 0$$

- They should be non-negative (or a penalty may turn into a reward):

$$\lambda \geq 0$$

- They should be 0 for all satisfied constraints (or  $\mathcal{L}$  would be "inflated")

$$\lambda \odot g(x) = 0$$

Additionally,  $x$  should be feasible, i.e.  $g(x) \leq 0$

## ...And KKT Conditions

If certain constraint qualifications apply, these are **necessary conditions**

If a point  $\mathbf{x}$  is a local optimum, then we have:

$$\nabla_{\mathbf{x}} \mathcal{L}(\mathbf{x}, \lambda) = 0 \quad (\text{null gradient})$$

$$\lambda \geq 0 \quad (\text{dual feasibility})$$

$$\lambda \odot g(\mathbf{x}) = 0 \quad (\text{complementary slackness})$$

$$g(\mathbf{x}) \leq 0 \quad (\text{primal feasibility})$$

They are known as Karush-Kuhn-Tucker (KKT) first order optimality conditions

### Some comments:

- If  $f(\mathbf{x})$  and  $g(\mathbf{x})$  are convex, the KKT conditions **are also sufficient**
- Equality constraints are equivalent to  $g(\mathbf{x}) \leq 0$  and  $-g(\mathbf{x}) \leq 0$
- ...Which can be manipulated to obtain (slightly) simpler formulas

# How to Use the KKT Conditions

We can use the KKT conditions to **constrain**  $x$  to be an optimum

- This is useful in bi-level optimization, i.e.:

$$\operatorname{argmax}_y \left\{ f(z) \mid z = \operatorname{argmin}_{x \in X} g(x, y) \right\}$$

- If  $X$  and  $g$  are convex, we can use the KKT conditions as constraints
- ...And replace the optimization step  $\operatorname{argmin}_{x \in X} g(x, y)$
- Typically, this is useful only if the conditions reduce to a simple form

We can use the KKT conditions to **check** whether  $x$  is a local optimum

...Assuming we are in convex optimization and constraint qualifications are met

- If we fix  $x$ , then the KKT conditions reduce to a linear system
- ...If we can solve it, then  $x$  is a local optimum
- ...And we have found the corresponding optimal multipliers

# How to Use the KKT Conditions in CG

As a by-product of the previous use case...

**If we know that  $x$  is an optimum, we can obtain the optimal  $\lambda$**

This is the application we care about

- If we have constraints in the master problem
- ...Rather than searching for variables such that:

$$\frac{\partial}{\partial x_j} f(x) < 0$$

- ...We search instead for variables such that:

$$\frac{\partial}{\partial x_j} \mathcal{L}(x, \lambda) < 0$$

## CG for the Modified Path Formulation

The modified Path Formulation can be rewritten as:

$$\arg \min_x \left\{ \frac{1}{2} (\|Vx - v\|_2^2 + \|Ex - e\|_2^2) + \alpha x \mid Vx \geq \gamma v, x \geq 0 \right\}$$

From which we obtain:

$$\mathcal{L}(x, \lambda, \mu) = \frac{1}{2} (\|Vx - v\|_2^2 + \|Ex - e\|_2^2) + \alpha x + \lambda^T (\gamma v - Vx) - \mu^T x$$

And finally:

$$\frac{\partial}{\partial x_j} \mathcal{L}(x, \lambda, \mu) = \sum_{i=1}^{n_v} r_i^v V_{ij} + \sum_{k=1}^{n_e} r_k^e E_{kj} - \sum_{i=1}^{n_v} \lambda_i V_{ij} + \alpha - \mu_j$$

## CG for the Modified Path Formulation

Therefore, we can modify our pricing problem so that we minimize:

$$\sum_{i=1}^{n_v} (r_i^v - \lambda_i) V_{ij} + \sum_{k=1}^{n_e} r_k^e E_{kj} + \alpha - \mu_j$$

Whenever we include an **arc**  $k$  in the path we are constructing:

- We accumulate a gradient term equal to  $r_k^e$
- ...Exactly the same as before

Whenever we include a **node**  $i$  in the path we are constructing:

- We accumulate a gradient term equal to  $r_i^v - \lambda_i$
- I.e. we subtract the multiplier associated to the  $i$ -th min cover constraint

Then, for every path, we add  $\alpha$  and we subtract  $\mu_j$

**But how do we get the multipliers?**



## CG for the Modified Path Formulation

**Every  $\lambda_i$  is associated to a (min cover) constraint:**

...And we have **all of those** in our problem!

- Hence, we could compute  $\lambda$  for the current optimal solution
- In practice, the OSQP solver can **compute  $\lambda$  for us**

**Every  $\mu_j$  is associated to  $x_j \geq 0$  constraint:**

...And unfortunately we have those only for the paths in the pool

- However, we know that  $\mu_j \geq 0$  (by dual feasibility)
- ...And we would have  $\mu_j > 0$  only having  $x_j < 0$  was beneficial
- ...But we are looking for paths with exactly the opposite property
- Hence, we can just assume  **$\mu_j = 0$**  when generating new paths

## Full CG In Action

---

This is Column Generation as it was meant to be

# Obtaining the Duals

## We start by solving again the master problem

...But this time we retrieve the optimal (dual) multipliers

- They are the same weights  $\lambda$  used in the ADMM
- ...And can be obtain from the OSQP solution object


```
In [11]: mvc, alpha = 0.95, 1
          rflows_n3, rpaths_n3, nneg_duals3, mvc_duals3 = util.solve_path_selection_full(tug, node_counts_n,
                                                                                       alpha=alpha, verbose=0, min_vertex_cover=0)
          print('FLOW: PATH')
          util.print_solution(tug, rflows_n3, rpaths_n3, sort='descending', max_paths=6)
          sse = util.get_reconstruction_error(tug, rflows_n3, rpaths_n3, node_counts_n, arc_counts_n)
          print(f'RSSE: {np.sqrt(sse):.2f}')
```

```
FLOW: PATH
8.28: 2,3 > 3,3
4.36: 0,2 > 1,2 > 2,2 > 3,2
2.62: 3,3
2.45: 0,1 > 1,1 > 2,0 > 3,0
1.68: 1,0 > 2,0 > 3,0
1.68: 0,1 > 1,1 > 2,0 > 3,2
...
RSSE: 2.54
```

# Inspecting the Duals

Let's inspect the multipliers for the minimum cover constraints

```
In [12]: util.plot_dict({i:v for i, v in enumerate(mvc_duals3)}, figsize=figsize)
```




- Some values are 0 (when the constraint is satisfied with a slack)
- ...And some (unexpectedly, are negative)

# Inspecting the Duals

Let's inspect the multipliers for the minimum cover constraints

```
In [13]: util.plot_dict({i:v for i, v in enumerate(mvc_duals3)}, figsize=figsize)
```




- The reason is the constraint direction: we have  $Vx \geq \gamma \hat{v}$  and not  $-Vx \leq -\gamma \hat{v}$
- We could fix by switching the constraint direction

# Inspecting the Duals

Let's inspect the multipliers for the minimum cover constraints

```
In [14]: util.plot_dict({i:v for i, v in enumerate(mvc_duals3)}, figsize=figsize)
```



- ...Or by reworking the change through the KKT formulas
- In our case, when we include a node we add  $\lambda_i$  instead of subtracting it

# Checking the Pricing Solution

**Our pricing problem code can handle both the duals and the L1 weight**

We modify node residuals by adding the cover multipliers:

```
if cover_duals is not None:
    for i, v in enumerate(tug.vs):
        nk = v['time'], v['index_o']
        nres[nk] += cover_duals[i]
```

- This provides an **incentive** to select paths
- ...That traverse a node whose cover constraint is satisfied with a slack

And we add the constant  $\alpha$  to the final path weights:

```
spw = [v + alpha for v in spw]
```

- This a uniform **disincentive** so select paths

**The shortest paths problem is solved as usual**

# Running the CG Approach

We can now run the CG approach

```
In [15]: rflows_cg, rpaths_cg = util.trajectory_extraction_cg(tug, node_counts_n, arc_counts_n,
                                                         alpha=alpha, min_vertex_cover=mvc, max_iter=30,
                                                         verbose=1, max_paths_per_iter=10, solver='piqp')
sse = util.get_reconstruction_error(tug, rflows_cg, rpaths_cg, node_counts_n, arc_counts_n)
print(f'RSSE: {np.sqrt(sse):.2f}')
```

```
It.0, sse: 209.13, #paths: 26, new: 10
It.1, sse: 83.13, #paths: 34, new: 8
It.2, sse: 67.39, #paths: 41, new: 7
It.3, sse: 42.31, #paths: 45, new: 4
It.4, sse: 7.24, #paths: 47, new: 2
It.5, sse: 6.46, #paths: 49, new: 2
It.6, sse: 6.46, #paths: 49, new: 0
RSSE: 2.54
```

```
/Users/michelelombardi/Library/Caches/pypoetry/virtualenvs/mining-with-co-eC0Z9-8k-py3.11/
lib/python3.11/site-packages/scipy/sparse/_index.py:108: SparseEfficiencyWarning: Changing
the sparsity structure of a csc_matrix is expensive. lil_matrix is more efficient.
  self._set_intXint(row, col, x.flat[0])
```

- Indeed, we obtain the same RSSE as the approach using all paths

**We now know how to use CG with constraints in the master problem**



## Some References

CG is main studied in the context of Linear Optimization

- "Column Generation", by Lübbecke
- "Column Generation", by Desaulniers, Desrosiers, and Solomon