

# Constraints in the Subproblem

---

When troubles spring up like mushrooms

# User Habits

**What if we know something about the habits of our users?**

- E.g. we may know that they don't tend to spend a long time on a single page
- We could use this information to further regularize the problem

**Specifically, we could add a constraint in the subproblem**

I.e. by putting a limit on consecutive visits to the same node

- It seems simple enough, but in practice it's serious issue

**Why is that the case?**

# User Habits

**What if we know something about the habits of our users?**

- E.g. we may know that they don't tend to spend a long time on a single page
- We could use this information to further regularize the problem

**Specifically, we could add a constraint in the subproblem**

I.e. by putting a limit on consecutive visits to the same node

- It seems simple enough, but in practice it's serious issue

**Why is that the case?**

- Such a constraint violates a basic assumption in Dijkstra's method
- I.e. that all path information can be condensed into its length

**With the new constraint, our shortest path method no longer works**

# Walking the Line

## **With our shortest path approach, we were walking a fine line**

- The problem could be solved in polynomial time
- ...But even a small addition could make it NP-hard instead

## **With the new constraint, pricing becomes indeed NP-hard**

There is nothing we can do about that

- ...But perhaps we can use a better suited technique
- Something designed specifically for NP-hard, combinatorial problems
- ...With lots of messy constraints

# Walking the Line

## **With our shortest path approach, we were walking a fine line**

- The problem could be solved in polynomial time
- ...But even a small addition could make it NP-hard instead

## **With the new constraint, pricing becomes indeed NP-hard**

There is nothing we can do about that

- ...But perhaps we can use a better suited technique
- Something designed specifically for NP-hard, combinatorial problems
- ...With lots of messy constraints

## **For example, we could use Constraint Programming**

...In its more modern incarnation, **Lazy Clause Generation** (a.k.a. CP-SAT)

# Constraint Programming and Lazy Clause Generation

---

Very little is lazy about that

# Constraint Satisfaction Problems

CP is techniques designed to address Constraint Satisfaction Problems (CSPs):

$$\langle X, D, C \rangle$$

Where:

- $X$  is a set of decision variables
- $D$  is the set of their domains
- $C$  is a set of constraints
- $f$  is a cost function

**Almost any decision problem fits those definitions...**

...But in practice, a given CP solver provides

- A library of supported variables types
- A library of supported constraints

## ...And Constraint Optimization Problems

CP can handle Constraint Optimization Problems (COP):

$$\langle f, X, D, C \rangle$$

- Where  $f$  is a cost function

COPs are tackled as a **sequence of CSPs** via this scheme:

- best solution  $x^* = \perp$
- while true find a solution for  $\langle X, D, C \rangle$ 
  - If a solution  $x'$  is found:
    - $x^* = x'$
    - $C = C \cup \{f(x) < f(x')\}$  # We as for an improving solution
  - otherwise, break the loop

The solver state is **maintained between solutions** so as not to waste effort



# Variables and Constraints

## In terms of supported variables types

- All CP solver provide **integer** variables
- Some also provide **numeric**, **interval**, **set**, or **graph** variables

## In terms of supported constraints

- All CP solvers provide **equalities**, **inequalities**, \_ over **linear expressions**
  - $y = a^T \mathbf{x}, y \leq a^T \mathbf{x}$
- All CP solvers provide  $\neq$  constraints
  - $y \neq x$
- Most CP solvers provide **max** and **min** constraints
  - $y = \max(\mathbf{x}), y = \min(\mathbf{x})$
- Some CP solvers provide products and modulo constraints (over scalars)
  - $z = xy, y = x \mod a$

# Variables and Constraints

CP solver provide also **constraints with non-mathematical nature**

- E.g. logical constraints:

$$x \vee y, x \wedge y, x \Rightarrow y \dots$$

- E.g. a set of variables should take all different values:

$$\text{ALLDIFFERENT}(x)$$

- E.g. a set of variables should take/not take values from a table  $T$ :

$$\text{ALLOWED}(x, T) \quad \text{and} \quad \text{FORBIDDEN}(x, T)$$

- E.g. a set of activities with start times  $x$  and durations  $d$  should not overlap:

$$\text{NOOVERLAP}(x, d)$$

# Propagators

## CP solvers are **search based**

They maintain information about the variable domains in a **Domain Store**:

- The solver may store the domain bounds, i.e.  $x_i \in \{lb_i, \dots, ub_i\}$
- ...Or the individual allowed values, i.e.  $x_i \in \{v_0, v_1, \dots\}$
- Other representations are also possible

## Constraints are associated to algorithms called **propagators**

- A propagator takes as input the current variable domains
- ...And can **prune** (some) provably infeasible values

By doing so, we can dramatically reduce the size of the search space

## Propagators often rely on **structural patterns** to improve pruning

- E.g. ALLDIFFERENT( $x$ ) can prune more than  $x_i \neq x_j, \forall i \neq j$
- ...Since it can reason on multiple variables at the same time

# Propagators

Let's see an example for  $\text{ALLOWED}([x_0, x_1], T)$ , with  $T$  given by:

$x_0$	$x_1$
0	0
0	1
1	1

Let's that initially  $D_0 = \{0, 1\}$ , and  $D_1 = \{0, 1\}$

- If  $x_0$  loses the value 0
  - ...Then the **ALLOWED** propagator prunes 0 from  $D_1$
  - ...Because it no longer has a feasible support in  $D_0$
- If  $x_1$  loses the value 1
  - ...Then the **ALLOWED** propagator prunes 1 from  $D_0$
  - ...Because it no longer has a feasible support in  $D_1$

# Propagators and Lazy Clause Generation

In Lazy Clause Generation solvers, propagators have two additional tasks:

## 1) Whenever they prune a domain, they also **generate boolean literals**

- These correspond to the pruning operations
  - E.g. in our two example we would generate literals  $[x_0 \neq 0]$  and  $[x_1 \neq 1]$
- These literals represent **variables** associated to the state of a constraint
  - E.g.  $[x_0 \neq 0] = 1$  if  $0 \notin D_0$  and  $[x_0 \neq 0] = 0$  otherwise

## 2) Whenever they prune a domain, they also generate an **explanation**

- This is a logical **clause** representing the reasoning that led to pruning
  - E.g. in our first example we would generate  $[x_0 \neq 0] \Rightarrow [x_1 \neq 0]$
- These clauses are constraints on the literal variables
  - They function like normal constraints (except they are specifically tracked)

# Constraint Propagation

## Pruning can trigger the activation of other propagators

...In a process called **Constraint Propagation**

- After some activations, the process reaches a **fix point**
- If propagation causes a domain to become empty, we have a **conflict**
- ...In which case we need to backtrack

## As a consequence, propagators are called many times per search node

We can have millions of propagator calls in a solution process

- For this reason, they often run in **constant or low-degree polynomial** time
- ...And propagator are heavily optimized
  - E.g. the **ALLOWED** constraint is so efficient
  - ...That tables with  $> 100,000$  entry can be handled almost instantly
  - This is achieved by relying on incremental computation

# Constraint Propagation and Implication Graphs

In LCG solvers, constraint propagation generates an **implication graph**

This consists of the **literals**, connected by the generated **explanations**

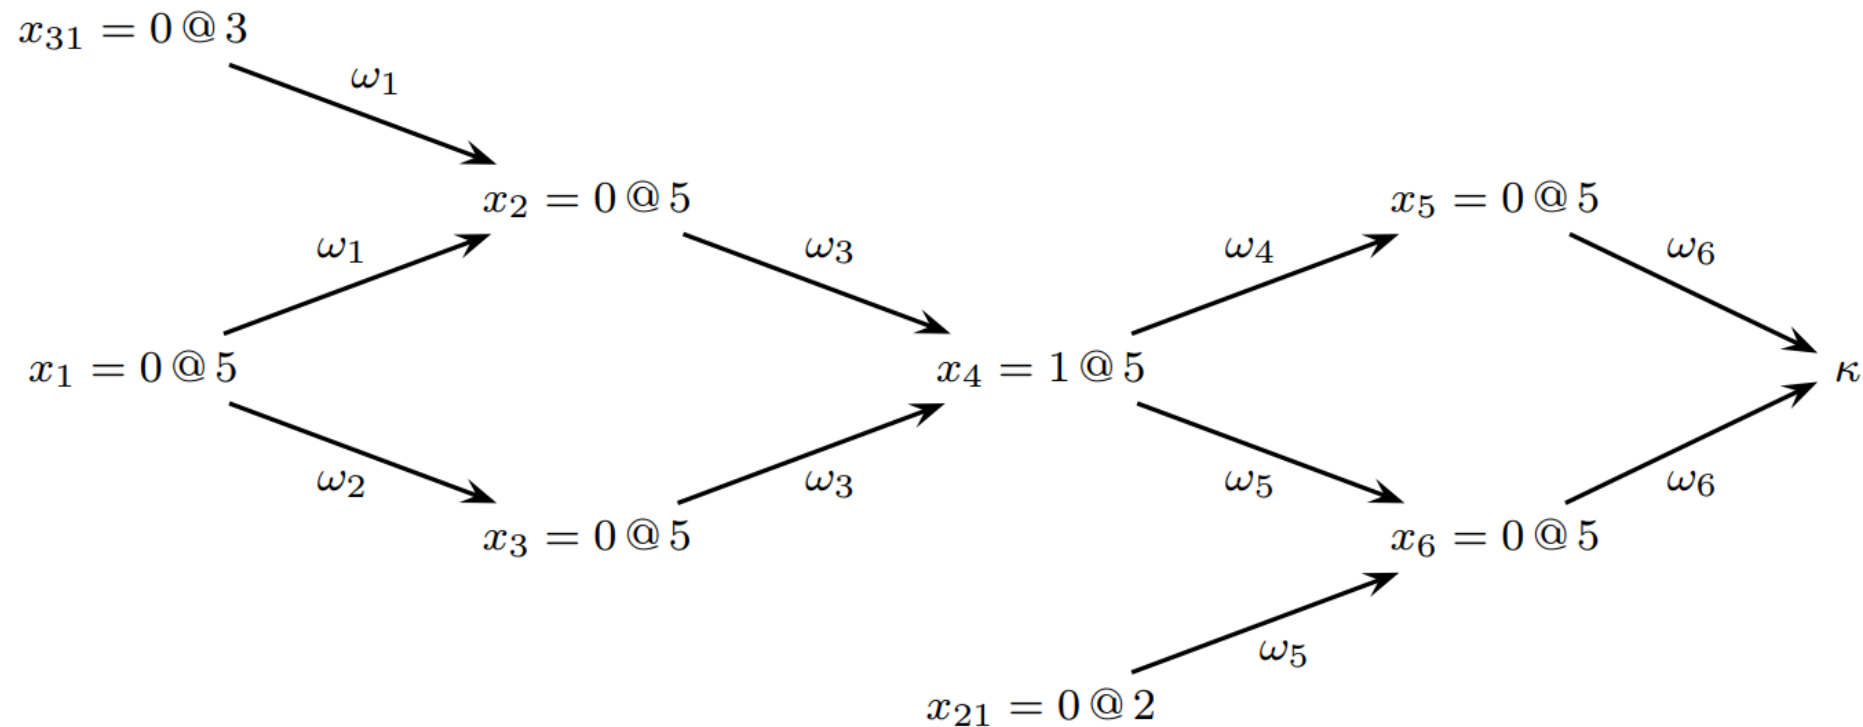


Image from [this book chapter](#)

- In the pictures,  **$x$**  variables correspond to boolean literals (i.e. constraint states)
- ...And each  **$\omega$**  is an explanation

# Conflict Driven Clause Learning

In LCG solvers, each search decision also generates a literal

E.g. if we assign 1 to  $x_0$ , we generate  $[x_0 = 1]$

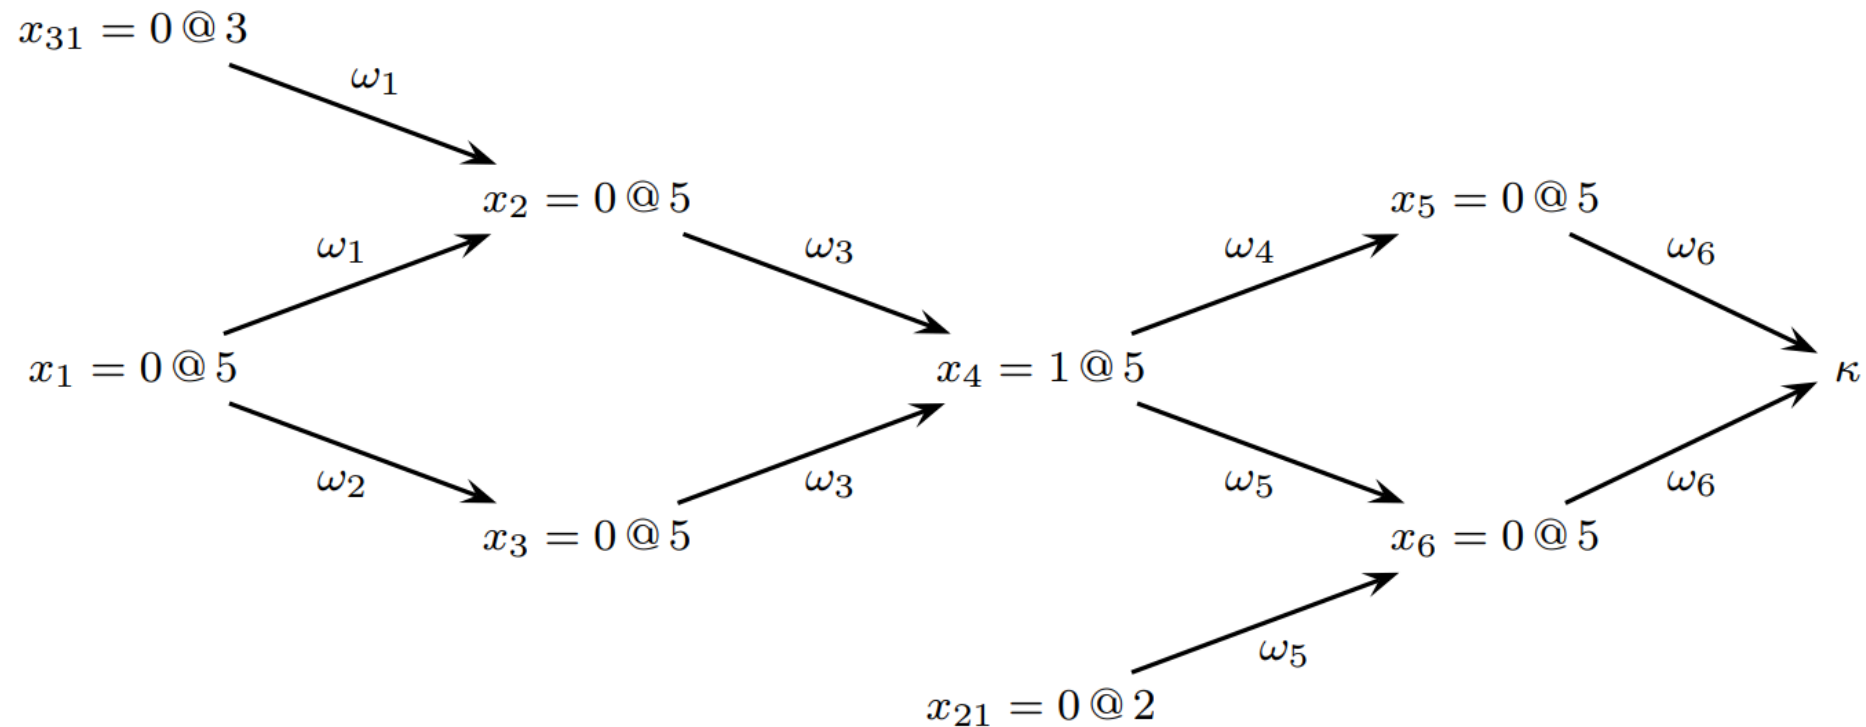


Image from [this book chapter](#)

Each decisions is associated to a **decision value**

- In the picture, they are the number after the @ symbol
- When we make a new decision, we increment the current decision value



# Conflict Driven Clause Learning

In LCG solvers, each search decision also generates a literal

E.g. if we assign 1 to  $x_0$ , we generate  $[x_0 = 1]$




Image from [this book chapter](#)

Literals generated by propagation are also labeled with a decision value

- ...But in this case there is **no increment**
- In the picture, many literals are associated to decision level 5

# Conflict Driven Clause Learning

In case of a conflict ( $\kappa$  in the figure), an LCG solver can **learn a constraint**

This technique is referred to as **Conflict Driven Clause Learning**

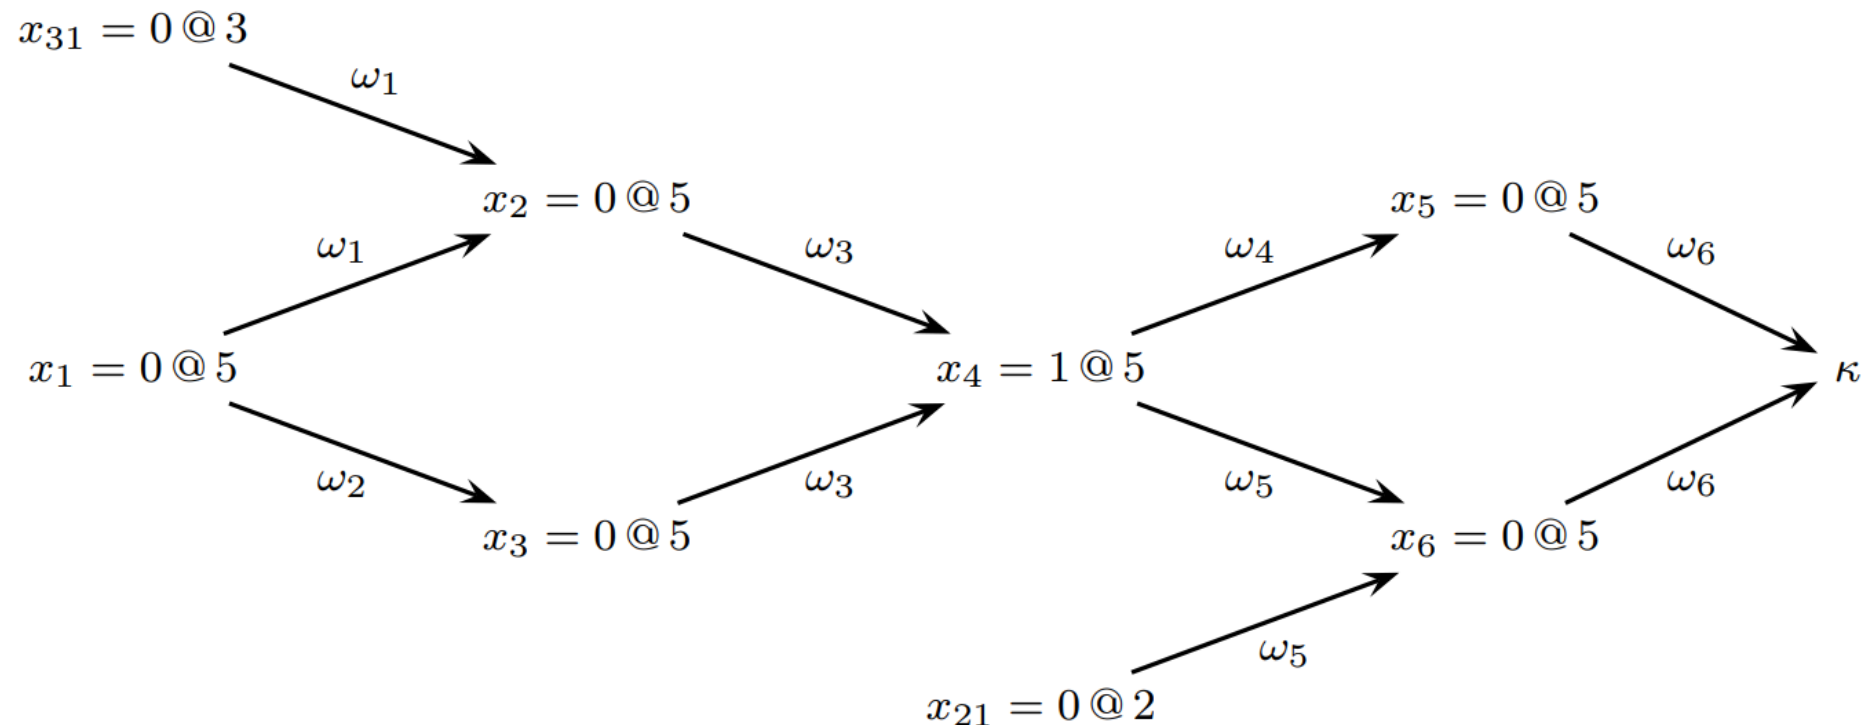


Image from [this book chapter](#)

The idea is to identify **which decisions (literals) are to blame** for the conflict

- First, we identify all literal with the **same decision value** as the conflict
- The earliest one ( $x_1 = 0@5$  in the figure) always corresponds to a decision

# Conflict Driven Clause Learning

In case of a conflict ( $\kappa$  in the figure), an LCG solver can **learn a constraint**

This technique is referred to as **Conflict Driven Clause Learning**

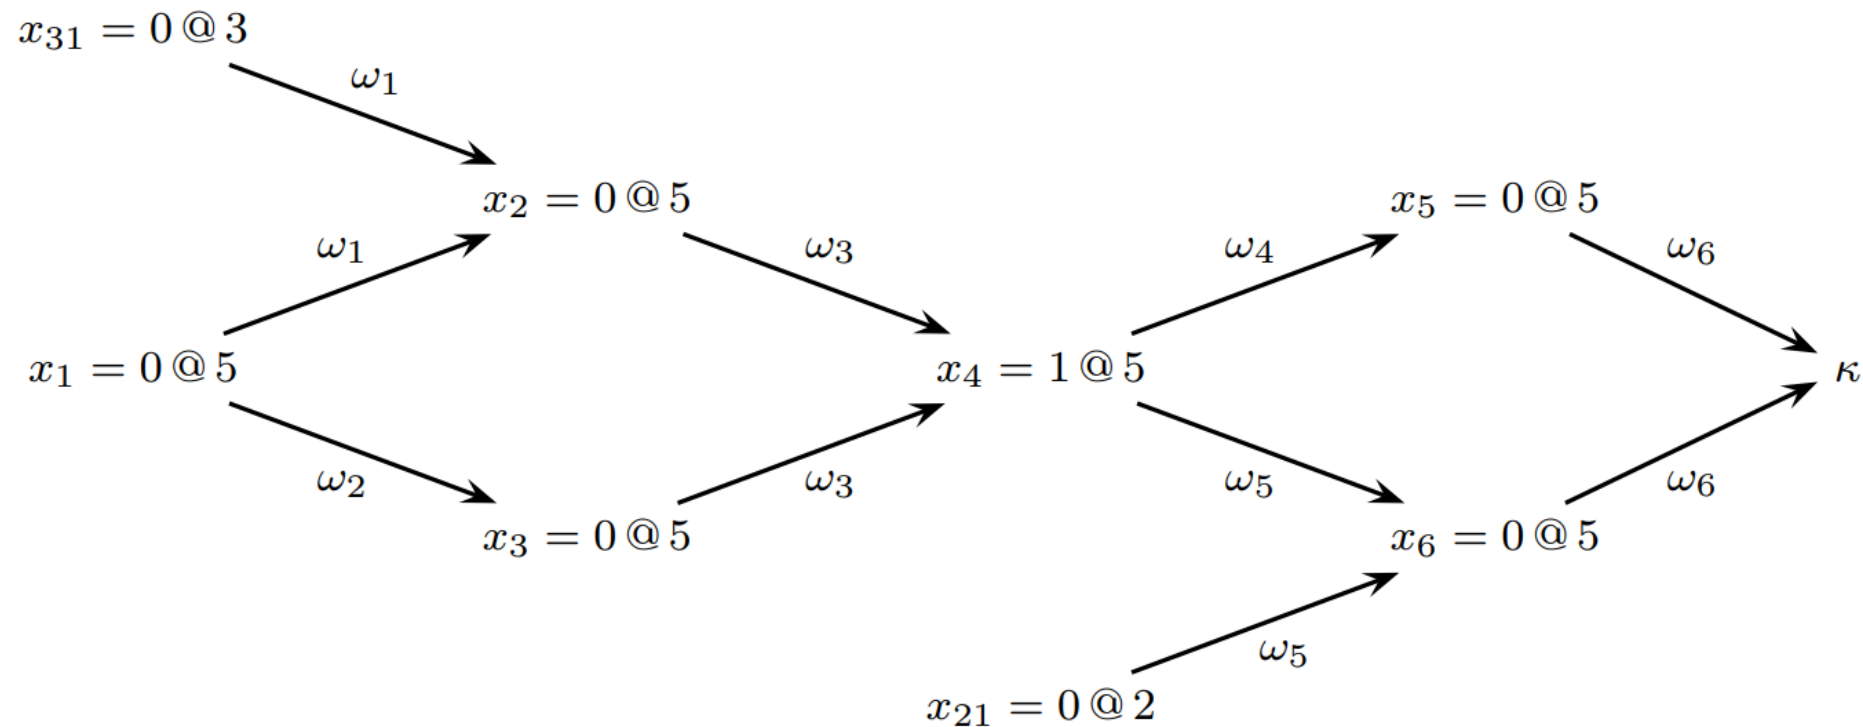


Image from [this book chapter](#)

Without this literal, the conflict would not arise

- ...Therefore it will appear in the learned constraint

# Conflict Driven Clause Learning

In case of a conflict ( $\kappa$  in the figure), an LCG solver can **learn a constraint**

This technique is referred to as **Conflict Driven Clause Learning**

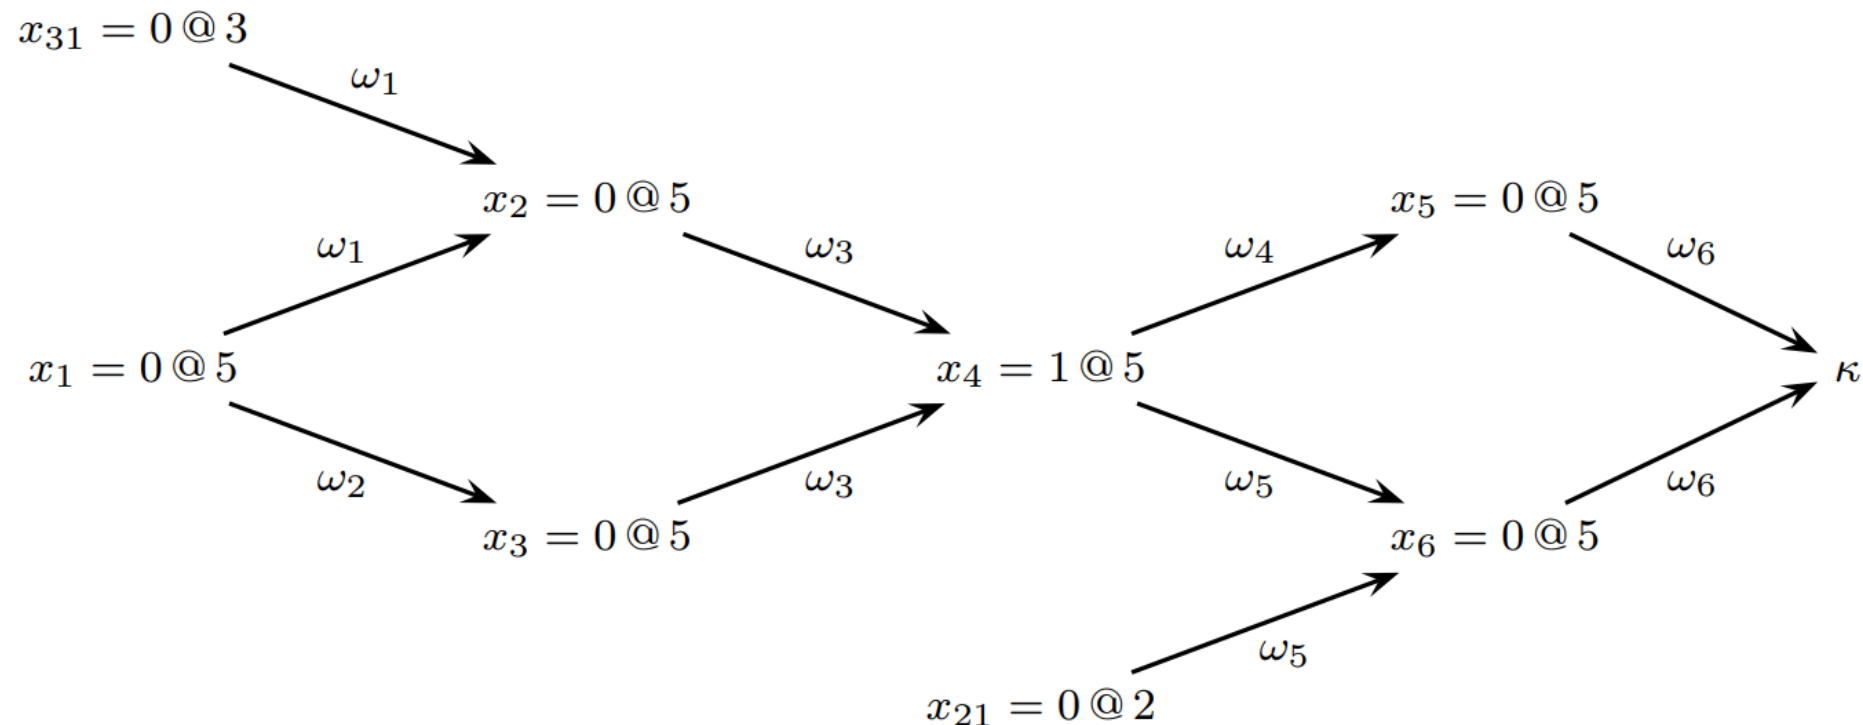


Image from [this book chapter](#)

To this, we add all literals with decision level **lower** than the current one

- ...That are connected via explanation to literals in the current decision level
- In the figure, those would be  $x_{31} = 0@3$  and  $x_{21} = 0@2$

# Conflict Driven Clause Learning

In case of a conflict ( $\kappa$  in the figure), an LCG solver can **learn a constraint**

This technique is referred to as **Conflict Driven Clause Learning**

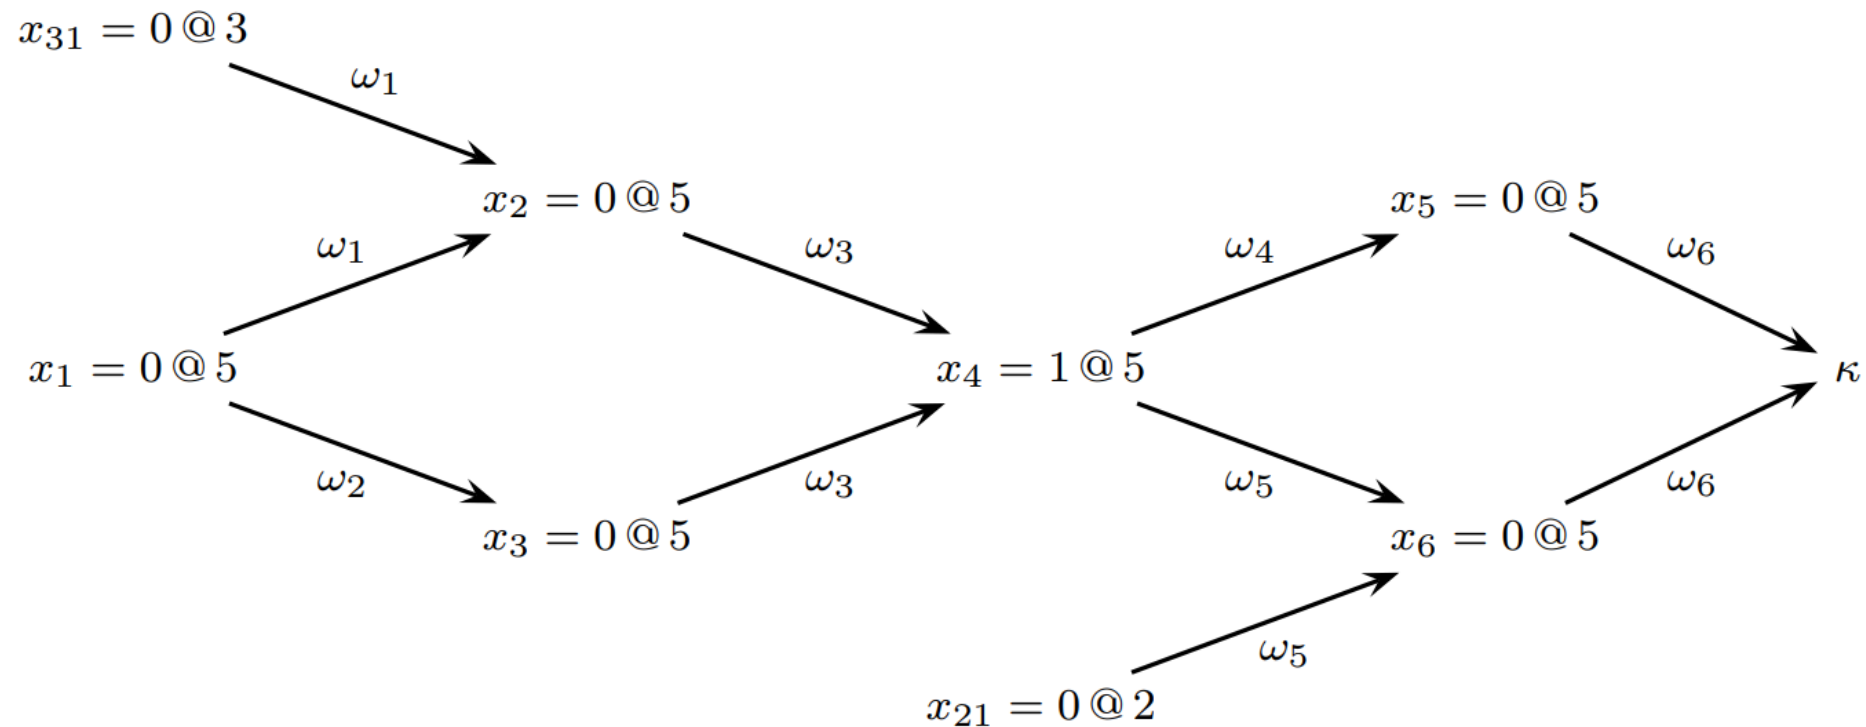


Image from [this book chapter](#)

If we want to avoid the conflict, **at least one** of these literals should be **false**

Therefore the clause we learn is:  $\neg[x_1 = 0] \vee \neg[x_{31} = 0] \vee \neg[x_{21} = 0]$

# Conflict Driven Clause Learning

The clause we learn is **globally valid**

...So that we can restart search and we will not make the same mistake again

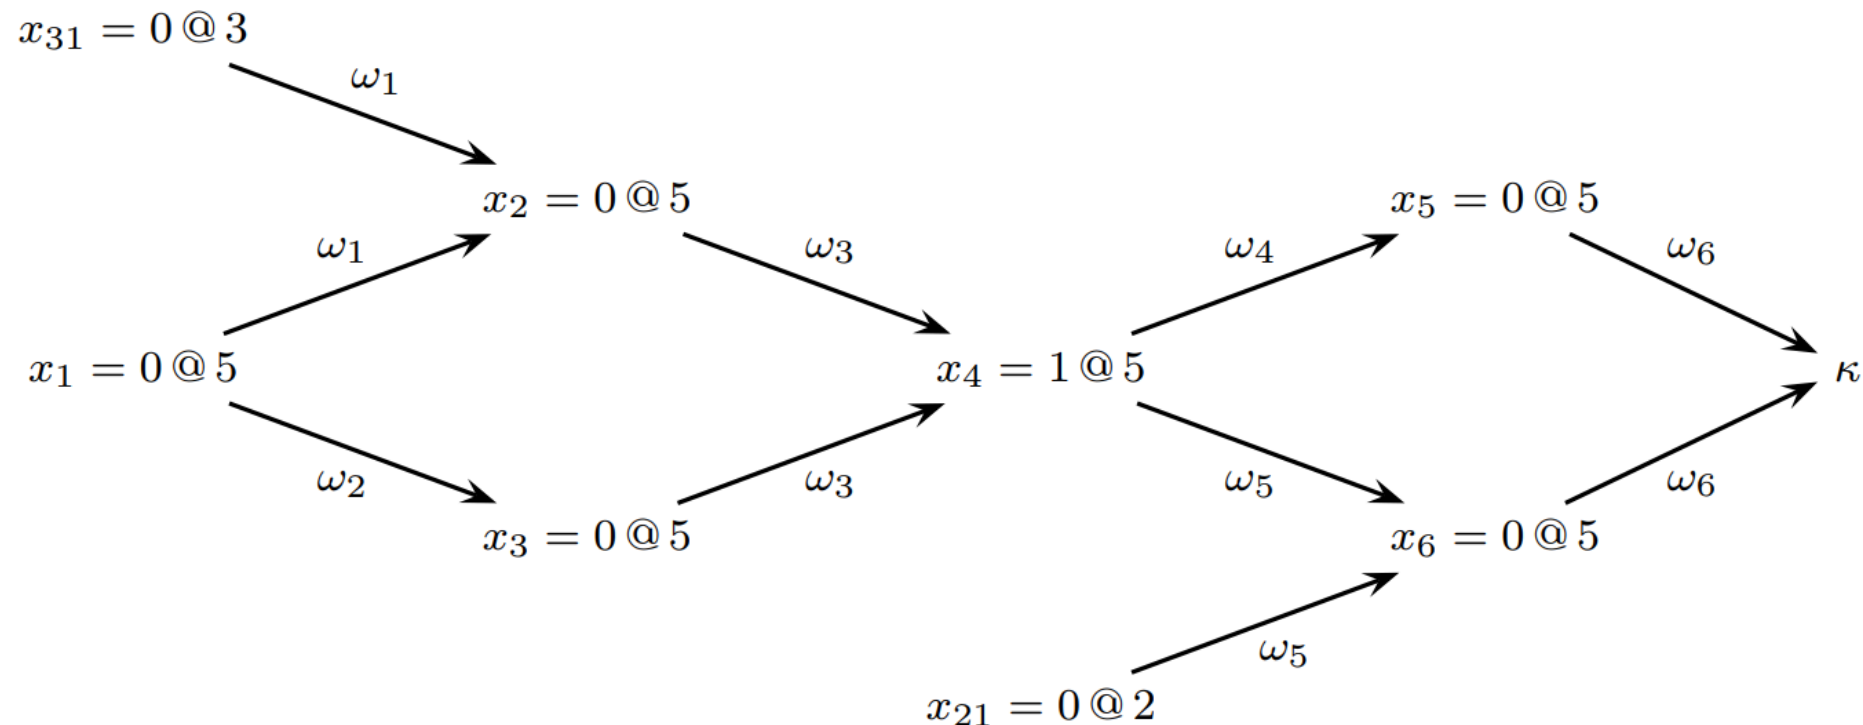


Image from [this book chapter](#)

In other words, we have a **complete** method that does **not** rely on tree search

- CDCL was invented for pure SAT solvers, and it is key to their efficiency
- In LCG we used it to obtain similarly strong benefits

## Some Considerations

**As usual, we have just scratched the surface for CP/LCG**

- You can find more information about classical CP [in this handbook](#)
- ...And for LCG the best starting point are the papers by [Peter Stuckey](#).

**Unlike MILP, CP does not rely on numerical optimization**

- Combinatorial constraints are first-class citizens
  - E.g. we have no big-Ms here!
- It tends to work best for problems with many combinatorial elements

**Unlike MILP, CP lack a global bounding method**

- There is no LP relaxation, and propagation works at a **local** level
- CDCL goes a long way towards countering this issue
- ...But sometimes the lack of a global bound leads to weaker performance

## Some References

If you are interested, you might want to check:

- "Handbook of Constraint Programming", by Rossi, Van Beek, Walsh
- "Satisfiability Modulo Theories", by Barrit, Sebastiani, Sanjit, Seshia, and Tinelli  
(a chapter from the "Handbook of Satisfiability")